

TF-IDF Vectorization

Instructions Using either Python or R for this assignment, program your own function to take a set of text files and perform TF-IDF vectorization on them.

Your work can be programmed in the .R, .rmd, .py, or .ipynb format.

Please compare the similarity of the files as a result of this.

Strategy

1. Load all text documents and vectorize == aggregate as one *Master list* of words
2. For each document: Count how often the words in the master list appear == Term Frequency (**TF**) vectors for each document
3. For each word: Calculate the Inverse Document Frequency (**IDF**) to create an IDF vector as follows:

$$\log(\# \text{ of total documents}) / [\log(\# \text{ documents using specific word} + 1)]$$

File similarity from these vectors:

- Goal of IDF: if a word appears in almost every document (eg. "the", "and") the IDF evaluates to ~0.
- Dot product of TF (per document) and IDF (weights of word importance): outputs a scalar & is based on how closely the two align.
- These output vectors can be analyzed with clustering algorithms such as k-means to estimate how similar they are to one another.

```
In [1]: import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
from sklearn.neighbors import KNeighborsClassifier
```

Example

I will start by importing one text file and demonstrating the process that the function later will perform on all text files

Step 1: Import text file and create master list of unique words that appear

```
In [2]: master_list = []
with open("C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\Alice_in_Wonderland.txt") as f:
    for line in Alice:
        words = line.split()
        for word in range(0, len(words)):
            if words[word] not in master_list:
                master_list.append(words[word])
```

```
In [3]: #view the master list as needed
#master_list
```

```
In [4]: #Checking the master list before proceeding
unique_check = np.unique(master_list, return_counts=True) #return counts parameter = 1
print(unique_check)

unique_count = 0
not_unique = 0

for i in range(0, len(unique_check[1])): #index second position for the unique counts r
    if unique_check[1][i] == 1:
        unique_count += 1
    if unique_check[1][i] != 1:
        not_unique += 1

print(f'Number of unique entries in the master list: {unique_count}')
print(f'Number of DUPLICATE entries in the master list: {not_unique}')
```

```
(array(['(Alice', '(And,', '(As', ..., '“We', '“—found', '“’Tis'],
      dtype='<U46'), array([1, 1, 1, ..., 1, 1, 1], dtype=int64))
Number of unique entries in the master list: 5167
Number of DUPLICATE entries in the master list: 0
```

Step 2: Term-Frequency Vector

```
In [5]: # Use the master list to create Term-Frequency (TF) Vector

tf_Alice = np.zeros(shape=(len(master_list)), dtype=np.int32) #create placeholder array

with open("C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\Alice_in_Wonderland.txt") as f:
    for term in range(0, len(master_list)): #Loop through the master list
        Alice.seek(0) #return to the beginning of the document
        for line in Alice: #Loop through each line in the document (
            words = line.split()
            for word in range(0, len(words)): #Loop through each word in the document
                if words[word] == master_list[term]:
                    tf_Alice[term] += 1 #count each instance of the word, in the
```

```
In [6]: print(f'Length of TF vector for Alice in Wonderland: {len(tf_Alice)}')
print(f'Length of (example) master list {len(master_list)}')
print('')
print(f'The first 10 values in the TF vector for Alice in Wonderland are: \n{tf_Alice[0:10]}')
print(f'The corresponding words from the master list are: \n{master_list[0:10]}')
```

Length of TF vector for Alice in Wonderland: 5167
Length of (example) master list 5167

The first 10 values in the TF vector for Alice in Wonderland are:

```
[ 12  1  1 1513  1 221 332  11 706  43]
```

The corresponding words from the master list are:

```
['CHAPTER', 'I.', 'Down', 'the', 'Rabbit-Hole', 'Alice', 'was', 'beginning', 'to', 'get']
```

IDF Vector, dot product & Clustering

Calculated later when we have several documents to compare to each other.

Function

```
In [7]: #function will accept list of filenames
file_Alice = "C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\Alice_in_Wor
file_Two_Cities = "C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\A_Tale_c
file_Frankenstein = "C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\Frank
file_Great_Gatsby = "C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\The_Gr
file_Pride_and_Prejudice = "C:\\Users\\Cal\\Desktop\\DSCI 508 Machine Learning\\Week 7\\

In [8]: master_list = []

#this function combines Step 1 and Step 2 above, and then also creates an IDF vector f
def tf_idf_vector_func(files_vector): #accepts a list of filenames
    hold_tf_vectors = []
    hold_idf_vector = []
    hold_dot_products = []
    #Create master list of all words that appear in all documents (Similar to step 1 a
    for filename in range(0,len(files_vector)):
        with open(files_vector[filename], 'r', encoding='utf-8') as file:
            for line in file:
                words = line.split()
                for word in range(0,len(words)):
                    if words[word] not in master_list:
                        master_list.append(words[word])
    #Create tf vectors now that the master list is created (Similar to step 2 above)
    for filename in range(0,len(files_vector)):
        hold_tf_vectors.append(np.zeros(shape=(len(master_list)),dtype=np.int32)) #hol
        with open(files_vector[filename], 'r', encoding='utf-8') as file:
            for term in range(0,len(master_list)): #Loop through the master list
                file.seek(0) #return to the beginning of the document
                for line in file: #Loop through each line in the document
                    words = line.split()
                    for word in range(0,len(words)): #Loop through each word in the document
                        if words[word] == master_list[term]:
                            hold_tf_vectors[filename][term] += 1 #count each occurrence
    #Create vector for IDF
    for word in range(0,len(master_list)): #IDF vectors are per word, Loop through each word
        word_count = 0 #Count how often each word appears across all documents
        for vector in range(0,len(hold_tf_vectors)): #By looping over the tf vectors
            if hold_tf_vectors[vector][word] > 0:
                word_count += 1
        idf_value = np.log(len(files_vector))/np.log(word_count+1) #IDF = (# total documents) / (# documents containing word)
        hold_idf_vector.append(idf_value) #store IDF value in vector, same length as master list
    #Dot products of TF and IDF vectors
```

```

for tf_vector in range(0,len(hold_tf_vectors)):
    dot_product = np.dot(hold_tf_vectors[tf_vector],hold_idf_vector)
    hold_dot_products.append(dot_product)
return hold_tf_vectors, hold_idf_vector, hold_dot_products

```

In [9]: `vectorize_books = tf_idf_vector_func([file_Alice, file_Two_Cities, file_Frankenstein,`

Examine resulting vectors

In [10]: `#idf vector`
`print(f'The shape of the IDF vector is: {np.shape(vectorize_books[1])}')`
`print(f'The shape of the master list is: {np.shape(master_list)}')`
`print('')`
`#print(f'The IDF vector is: \n {vectorize_books[1]} ')`

The shape of the IDF vector is: (37357,)
The shape of the master list is: (37357,)

As expected, the IDF array holds only a single vector which has a value for each word index.
These numbers have a lot of floating point decimals as the result of division of logarithms.

In [11]: `#tf vectors`
`tf_Alice = vectorize_books[0][0]`
`tf_Two_Cities = vectorize_books[0][1]`
`tf_Frankenstein = vectorize_books[0][2]`
`tf_Great_Gatsby = vectorize_books[0][3]`
`tf_Pride_and_Prejudice = vectorize_books[0][4]`

`print(f'The first five words of the master list are: \n {master_list[0:5]}')`
`tf_array = [tf_Alice,tf_Two_Cities,tf_Frankenstein,tf_Great_Gatsby,tf_Pride_and_Prejudice]`

`for i in range(0,len(tf_array)):`
 `print(np.shape(tf_array[i]))`
 `print(tf_array[i][0:5])`

The first five words of the master list are:
['CHAPTER', 'I.', 'Down', 'the', 'Rabbit-Hole']
(37357,)
[12 1 1 1513 1]
(37357,)
[44 10 0 7347 0]
(37357,)
[0 1 0 3897 0]
(37357,)
[0 0 0 2201 0]
(37357,)
[59 0 0 4069 0]

As expected, all are the same length as the master list, and are whole numbers since they are integer counts of the words.

And also as expected, "the" is a very common word!

In [12]: `#dot products`
`#vectorize_books[2]`

```
for dots in range(0, len(vectorize_books[2])):
    print(vectorize_books[2][dots])
```

```
29327.867030090445
152801.7373287466
81668.23575407494
54957.31596921146
132854.09055423026
```

The dot product of two vectors outputs a scalar. **I wonder if element-wise matrix multiplication is actually more practical since it outputs a vector?**

```
In [13]: #element-wise multiplication rather than dot product

hadamard_products = []

for tf_vector in range(0, len(tf_array)):
    hadamard_prod = np.multiply(tf_array[tf_vector], vectorize_books[1])
    print(hadamard_prod[0:5])
    hadamard_products.append(hadamard_prod)
```

```
[1.39315686e+01 1.16096405e+00 2.32192809e+00 1.35904378e+03
 2.32192809e+00]
[ 51.08241809  11.60964047   0.          6599.40161932   0.          ]
[0.00000000e+00 1.16096405e+00 0.00000000e+00 3.50045843e+03
 0.00000000e+00]
[   0.          0.          0.          1977.03592815   0.          ]
[ 68.4968788   0.          0.          3654.95647053   0.          ]
```

Since these outputs are vectors, this is more suited to a clustering algorithm than scalars are.

Perform clustering on vectors & estimate similarity

```
In [14]: km_classifier = KMeans(n_clusters=2, random_state=0, n_init=10)
k_means_books = km_classifier.fit(X=hadamard_products)

print(f' The cluster assignment is as follows: {k_means_books.labels_}')

#View clusters if desired (very high dimensionality)
#k_means_books.cluster_centers_
```

```
The cluster assignment is as follows: [0 1 0 0 1]
```

```
C:\Users\Cal\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
  warnings.warn(
```

The assigned clusters are [0 1 0 0 1]. From this method, this conveys that books 1,3 and 4 are similar, while books 2, and 5 are similar.

Books 1,3,4 are:

- Alice in Wonderland
- Frankenstein

- The Great Gatsby

Books 2 and 5 are:

- A Tale of Two Cities
- Pride and Prejudice

I think it's understandable that A Tale of Two Cities and Pride and Prejudice are rather alike!